

Healthcare.capstone

January 19, 2024

DATASCIENCE CAPSTONE PROJECT

HEALTHCARE BY SUJATA KINHEKAR

```
[1]: #IMPORT LIBRARIES
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: data = pd.read_csv('health care diabetes.csv')
```

```
[3]: data.head()
```

```
[3]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0             6     148             72             35         0  33.6
1             1      85             66             29         0  26.6
2             8     183             64              0         0  23.3
3             1      89             66             23        94  28.1
4             0     137             40             35       168  43.1

      DiabetesPedigreeFunction  Age  Outcome
0                0.627     50         1
1                0.351     31         0
2                0.672     32         1
3                0.167     21         0
4                2.288     33         1
```

```
[4]: data.isnull().any()
```

```
[4]: Pregnancies      False
Glucose              False
BloodPressure        False
SkinThickness        False
Insulin              False
BMI                  False
DiabetesPedigreeFunction  False
Age                  False
Outcome              False
```

dtype: bool

```
[5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies           768 non-null    int64
1   Glucose               768 non-null    int64
2   BloodPressure         768 non-null    int64
3   SkinThickness         768 non-null    int64
4   Insulin               768 non-null    int64
5   BMI                   768 non-null    float64
6   DiabetesPedigreeFunction 768 non-null    float64
7   Age                   768 non-null    int64
8   Outcome               768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```


0.1 USING HISTOGRAMS

```
[6]: data['Glucose'].value_counts().head(7)
```

```
[6]: 100    17
      99     17
      129    14
      125    14
      111    14
      106    14
      95     13
      Name: Glucose, dtype: int64
```

```
[7]: plt.hist(data['Glucose'])
```

```
[7]: (array([ 5.,  0.,  4., 32., 156., 211., 163., 95., 56., 46.]),
      array([ 0. , 19.9, 39.8, 59.7, 79.6, 99.5, 119.4, 139.3, 159.2,
              179.1, 199. ]),
      <BarContainer object of 10 artists>)
```




```
[8]: data['BloodPressure'].value_counts().head(7)
```

```
[8]: 70    57
      74    52
      68    45
      78    45
      72    44
      64    43
      80    40
      Name: BloodPressure, dtype: int64
```

```
[9]: plt.hist(data['BloodPressure'])
```

```
[9]: (array([ 35.,   1.,   2.,  13., 107., 261., 243.,  87.,  14.,   5.]),
      array([ 0. , 12.2, 24.4, 36.6, 48.8, 61. , 73.2, 85.4, 97.6,
            109.8, 122. ]),
      <BarContainer object of 10 artists>)
```




```
[10]: data['SkinThickness'].value_counts().head(7)
```

```
[10]: 0      227
      32      31
      30      27
      27      23
      23      22
      33      20
      18      20
      Name: SkinThickness, dtype: int64
```

```
[11]: plt.hist(data['SkinThickness'])
```

```
[11]: (array([231., 107., 165., 175., 78., 9., 2., 0., 0., 1.]),
      array([ 0. , 9.9, 19.8, 29.7, 39.6, 49.5, 59.4, 69.3, 79.2, 89.1, 99. ]),
      <BarContainer object of 10 artists>)
```




```
[12]: data['Insulin'].value_counts().head(7)
```

```
[12]: 0      374
      105     11
      140      9
      130      9
      120      8
      100      7
       94      7
      Name: Insulin, dtype: int64
```

```
[13]: plt.hist(data['Insulin'])
```

```
[13]: (array([487., 155., 70., 30., 8., 9., 5., 1., 2., 1.]),
      array([ 0. , 84.6, 169.2, 253.8, 338.4, 423. , 507.6, 592.2, 676.8,
              761.4, 846. ]),
      <BarContainer object of 10 artists>)
```




```
[14]: data['BMI'].value_counts().head(7)
```

```
[14]: 32.0    13
      31.6    12
      31.2    12
      0.0    11
      33.3    10
      32.4    10
      32.8     9
      Name: BMI, dtype: int64
```

```
[15]: plt.hist(data['BMI'])
```

```
[15]: (array([ 11.,  0., 15., 156., 268., 224., 78., 12.,  3.,  1.]),
      array([ 0. ,  6.71, 13.42, 20.13, 26.84, 33.55, 40.26, 46.97, 53.68,
              60.39, 67.1 ]),
      <BarContainer object of 10 artists>)
```



```
[16]: data.describe().transpose()
```

```
[16]:
```

	count	mean	std	min	25%	\
Pregnancies	768.0	3.845052	3.369578	0.000	1.00000	
Glucose	768.0	120.894531	31.972618	0.000	99.00000	
BloodPressure	768.0	69.105469	19.355807	0.000	62.00000	
SkinThickness	768.0	20.536458	15.952218	0.000	0.00000	
Insulin	768.0	79.799479	115.244002	0.000	0.00000	
BMI	768.0	31.992578	7.884160	0.000	27.30000	
DiabetesPedigreeFunction	768.0	0.471876	0.331329	0.078	0.24375	
Age	768.0	33.240885	11.760232	21.000	24.00000	
Outcome	768.0	0.348958	0.476951	0.000	0.00000	

	50%	75%	max
Pregnancies	3.0000	6.00000	17.00
Glucose	117.0000	140.25000	199.00
BloodPressure	72.0000	80.00000	122.00
SkinThickness	23.0000	32.00000	99.00
Insulin	30.5000	127.25000	846.00
BMI	32.0000	36.60000	67.10
DiabetesPedigreeFunction	0.3725	0.62625	2.42
Age	29.0000	41.00000	81.00
Outcome	0.0000	1.00000	1.00

```
[17]: Positive=data[data['Outcome']==1]
      Positive.head(5)
```


```
[17]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI   \
0           6        148             72             35         0  33.6
2           8        183             64             0         0  23.3
4           0        137             40             35        168  43.1
6           3         78             50             32         88  31.0
8           2        197             70             45        543  30.5

      DiabetesPedigreeFunction  Age  Outcome
0                        0.627   50         1
2                        0.672   32         1
4                        2.288   33         1
6                        0.248   26         1
8                        0.158   53         1
```

0.2 USING HISTOGRAM STRENGTHEN THE MISSING VALUES

```
[18]: plt.hist(Positive['Glucose'],histtype='stepfilled',bins=20)
```

```
[18]: (array([ 2.,  0.,  0.,  0.,  0.,  0.,  0.,  1.,  4.,  9., 28., 26., 36.,
        27., 29., 22., 24., 21., 25., 14.]),
      array([ 0. ,  9.95, 19.9 , 29.85, 39.8 , 49.75, 59.7 , 69.65,
        79.6 , 89.55, 99.5 , 109.45, 119.4 , 129.35, 139.3 , 149.25,
        159.2 , 169.15, 179.1 , 189.05, 199.  ]),
      [<matplotlib.patches.Polygon at 0x7f9c984aad90>])
```




```
[20]: Positive['Glucose'].value_counts().head(7)
```



```
[20]: 125    7
      158    6
      128    6
      115    6
      129    6
      146    5
      162    5
      Name: Glucose, dtype: int64
```

```
[21]: plt.hist(Positive['BloodPressure'],histtype='stepfilled',bins=20)
```

```
[21]: (array([16.,  0.,  0.,  0.,  0.,  1.,  0.,  1.,  6.,  6., 19., 37., 56.,
        36., 41., 31.,  7.,  4.,  4.,  3.]),
      array([ 0. ,  5.7, 11.4, 17.1, 22.8, 28.5, 34.2, 39.9, 45.6,
        51.3, 57. , 62.7, 68.4, 74.1, 79.8, 85.5, 91.2, 96.9,
        102.6, 108.3, 114. ]),
      [<matplotlib.patches.Polygon at 0x7f9c982b8fd0>])
```




```
[23]: Positive['BloodPressure'].value_counts().head(7)
```

```
[23]: 70    23
      76    18
      78    17
      74    17
      72    16
      0     16
```

```
82    13
Name: BloodPressure, dtype: int64
```

```
[24]: plt.hist(Positive['SkinThickness'],histtype='stepfilled',bins=20)
```

```
[24]: (array([88.,  1.,  4., 10., 18., 30., 41., 34., 23., 15.,  1.,  1.,  1.,
          0.,  0.,  0.,  0.,  0.,  0.,  1.]),
      array([ 0. ,  4.95,  9.9 , 14.85, 19.8 , 24.75, 29.7 , 34.65, 39.6 ,
          44.55, 49.5 , 54.45, 59.4 , 64.35, 69.3 , 74.25, 79.2 , 84.15,
          89.1 , 94.05, 99.  ]),
      [<matplotlib.patches.Polygon at 0x7f9c98238a90>])
```




```
[26]: Positive['SkinThickness'].value_counts().head(7)
```

```
[26]: 0      88
      32     14
      33      9
      30      9
      39      8
      35      8
      36      8
      Name: SkinThickness, dtype: int64
```

```
[27]: plt.hist(Positive['Insulin'],histtype='stepfilled',bins=20)
```

```
[27]: (array([141.,  6., 23., 33., 24., 12.,  7.,  7.,  2.,  1.,  1.,
            5.,  3.,  1.,  1.,  0.,  0.,  0.,  0.,  1.]),
      array([ 0. , 42.3, 84.6, 126.9, 169.2, 211.5, 253.8, 296.1, 338.4,
            380.7, 423. , 465.3, 507.6, 549.9, 592.2, 634.5, 676.8, 719.1,
            761.4, 803.7, 846. ]),
      [<matplotlib.patches.Polygon at 0x7f9c981b4a50>])
```




```
[28]: Positive['Insulin'].value_counts().head(7)
```

```
[28]: 0      138
      130      6
      180      4
      156      3
      175      3
      194      2
      125      2
      Name: Insulin, dtype: int64
```

```
[29]: plt.hist(Positive['BMI'],histtype='stepfilled',bins=20)
```

```
[29]: (array([ 2.,  0.,  0.,  0.,  0.,  0.,  3., 13., 38., 61., 61., 36., 27.,
            14.,  7.,  3.,  1.,  1.,  0.,  1.]),
      array([ 0.   ,  3.355,  6.71 , 10.065, 13.42 , 16.775, 20.13 , 23.485,
            26.84 , 30.195, 33.55 , 36.905, 40.26 , 43.615, 46.97 , 50.325,
            53.68 , 57.035, 60.39 , 63.745, 67.1  ]),
      [<matplotlib.patches.Polygon at 0x7f9c981256d0>])
```



```
[30]: Positive['BMI'].value_counts().head(7)
```

```
[30]: 32.9      8
      31.6      7
      33.3      6
      30.5      5
      32.0      5
      31.2      5
      32.4      4
      Name: BMI, dtype: int64
```

```
[31]: Positive.describe().transpose()
```

```
[31]:
```

	count	mean	std	min	25%	\
Pregnancies	268.0	4.865672	3.741239	0.000	1.7500	
Glucose	268.0	141.257463	31.939622	0.000	119.0000	
BloodPressure	268.0	70.824627	21.491812	0.000	66.0000	
SkinThickness	268.0	22.164179	17.679711	0.000	0.0000	
Insulin	268.0	100.335821	138.689125	0.000	0.0000	
BMI	268.0	35.142537	7.262967	0.000	30.8000	
DiabetesPedigreeFunction	268.0	0.550500	0.372354	0.088	0.2625	
Age	268.0	37.067164	10.968254	21.000	28.0000	
Outcome	268.0	1.000000	0.000000	1.000	1.0000	


	50%	75%	max
Pregnancies	4.000	8.000	17.00

Glucose	140.000	167.000	199.00
BloodPressure	74.000	82.000	114.00
SkinThickness	27.000	36.000	99.00
Insulin	0.000	167.250	846.00
BMI	34.250	38.775	67.10
DiabetesPedigreeFunction	0.449	0.728	2.42
Age	36.000	44.000	70.00
Outcome	1.000	1.000	1.00


1 SCATTERPLOT

```
[32]: BloodPressure = Positive['BloodPressure']
      Glucose=Positive['Glucose']
      SkinThickness=Positive['SkinThickness']
      Insulin=Positive['Insulin']
```


```
[33]: plt.scatter(BloodPressure,Glucose,color=['blue'])
      plt.xlabel('BloodPressure')
      plt.ylabel('Glucose')
      plt.title('Glucose & BloodPressure')
      plt.show()
```




```
[34]: g=sns.scatterplot(x='Glucose',y='BloodPressure',hue='Outcome',data=data);
```



```
[35]: B=sns.scatterplot(x='BMI', y='Insulin', hue='Outcome', data=data);
```



```
[36]: g=sns.scatterplot(x='SkinThickness',y='Insulin',hue='Outcome', data=data);
```



1.1 CORRELATIONANALYSIS

```
[37]: data.corr()
```

```
[37]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	\
Pregnancies	1.000000	0.129459	0.141282	-0.081672	
Glucose	0.129459	1.000000	0.152590	0.057328	
BloodPressure	0.141282	0.152590	1.000000	0.207371	
SkinThickness	-0.081672	0.057328	0.207371	1.000000	
Insulin	-0.073535	0.331357	0.088933	0.436783	
BMI	0.017683	0.221071	0.281805	0.392573	
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928	
Age	0.544341	0.263514	0.239528	-0.113970	
Outcome	0.221898	0.466581	0.065068	0.074752	

	Insulin	BMI	DiabetesPedigreeFunction	\
Pregnancies	-0.073535	0.017683	-0.033523	
Glucose	0.331357	0.221071	0.137337	
BloodPressure	0.088933	0.281805	0.041265	
SkinThickness	0.436783	0.392573	0.183928	
Insulin	1.000000	0.197859	0.185071	
BMI	0.197859	1.000000	0.140647	
DiabetesPedigreeFunction	0.185071	0.140647	1.000000	


Age	-0.042163	0.036242	0.033561
Outcome	0.130548	0.292695	0.173844

	Age	Outcome
Pregnancies	0.544341	0.221898
Glucose	0.263514	0.466581
BloodPressure	0.239528	0.065068
SkinThickness	-0.113970	0.074752
Insulin	-0.042163	0.130548
BMI	0.036242	0.292695
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000

2 HEATMAPS


```
[38]: sns.heatmap(data.corr())
```

```
[38]: <AxesSubplot:>
```




```
[43]: plt.subplots(figsize=(8,8))
      sns.heatmap(data.corr(),annot=True,cmap='viridis')
```

[43]: <AxesSubplot:>



```
[45]: plt.subplots(figsize=(8,8))
      sns.heatmap(data.corr(), annot=True)
```

[45]: <AxesSubplot:>



WEEK 2: DATA MODELING

```
[46]: data.head()
```

```
[46]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
0           6     148           72           35         0  33.6
1           1      85           66           29         0  26.6
2           8     183           64            0         0  23.3
3           1      89           66           23        94  28.1
4           0     137           40           35       168  43.1

   DiabetesPedigreeFunction  Age  Outcome
```

0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
[47]: #TrainTestSplit
features=data.iloc[:,[0,1,2,3,4,5,6,7,]].values
label=data.iloc[:,8].values
```

```
[50]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(features,label,test_size=0.
↪2,random_state=10)
```

```
[52]: from sklearn.linear_model import LogisticRegression
model=LogisticRegression()
model.fit(X_train,y_train)
```

```
[52]: LogisticRegression()
```

```
[53]: print(model.score(X_train,y_train))
print(model.score(X_test,y_test))
```

```
0.7719869706840391
0.7662337662337663
```

```
[56]: from sklearn.metrics import confusion_matrix
cm=confusion_matrix(label,model.predict(features))
cm
```

```
[56]: array([[446,  54],
          [122, 146]])
```

```
[57]: from sklearn.metrics import classification_report
print(classification_report(label, model.predict(features)))
```

	precision	recall	f1-score	support
0	0.79	0.89	0.84	500
1	0.73	0.54	0.62	268
accuracy			0.77	768
macro avg	0.76	0.72	0.73	768
weighted avg	0.77	0.77	0.76	768

```
[61]: #PREPARING ROC CURVE(RECEIVER OPERATING CHARACTERISTICS CURVE)
```


```
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score

#PredictProbabilities
probs=model.predict_proba(features)

#Keep probabilities for the positive outcome only
probs=probs[:,1]
#Calculate AUC
auc=roc_auc_score(label,probs)
print('AUC: %.3f' % auc)
#calculate roc curve
fpr,tpr,thresholds=roc_curve(label,probs)
#plotno skill
plt.plot([0,1],[0,1],linestyle='--')
#plottheroccurveforthemodel
plt.plot(fpr, tpr, marker='.')
```

AUC: 0.837

```
[61]: [<matplotlib.lines.Line2D at 0x7f9c949e7150>]
```



```
[63]: #Applying decision tree classifier
from sklearn.tree import DecisionTreeClassifier
model3=DecisionTreeClassifier(max_depth=5)
```

```
model3.fit(X_train,y_train)
```

[63]: DecisionTreeClassifier(max_depth=5)

```
[64]: model3.score(X_train,y_train)
```

[64]: 0.8289902280130294

```
[65]: model3.score(X_test,y_test)
```

[65]: 0.7662337662337663

```
[66]: #APPLYING RANDOM FOREST  
from sklearn.ensemble import RandomForestClassifier  
model4=RandomForestClassifier(n_estimators=1)  
model4.fit(X_train,y_train)
```

[66]: RandomForestClassifier(n_estimators=1)

```
[68]: model4.score(X_train,y_train)
```

[68]: 0.8648208469055375

```
[69]: model4.score(X_test,y_test)
```

[69]: 0.7077922077922078

```
[70]: #Support Vector Classifier  
from sklearn.svm import SVC  
model5=SVC(kernel='rbf',gamma='auto')  
model5.fit(X_train,y_train)
```

[70]: SVC(gamma='auto')

```
[73]: model5.score(X_test,y_test),model5.score(X_train,y_train)
```

[73]: (0.6168831168831169, 1.0)

```
[74]: #Applying K-NN  
from sklearn.neighbors import KNeighborsClassifier  
model2=KNeighborsClassifier(n_neighbors=7, metric='minkowski',p=2)  
model2.fit(X_train,y_train)
```

[74]: KNeighborsClassifier(n_neighbors=7)

```
[75]: #Preparing ROC Curve(Receiver Operating Characteristics Curve)  
from sklearn.metrics import roc_curve
```

```

from sklearn.metrics import roc_auc_score

#predict probabilities
probs=model2.predict_proba(features)
#Keep probabilities for the positive outcome only
probs=probs[:,1]
#calculate AUC
auc=roc_auc_score(label,probs)
print('AUC: %.3F' % auc)
#Calculate roc curve
fpr, tpr, thresholds=roc_curve(label,probs)
print("TruePositiveRate-{}, FalsePositiveRate-{}Threshold-{}".
      format(tpr, fpr, thresholds))
#plotno skill
plt.plot([0,1],[0,1],linestyle='--')
#plottheroccurveforthemodel
plt.plot(fpr, tpr, marker='.')
plt.xlabel("FalsePositiveRate")
plt.ylabel("TruePositiveRate")

```


AUC: 0.836

```

TruePositiveRate-[0.      0.06716418 0.23880597 0.44776119 0.60074627
0.75373134
0.88059701 0.98507463 1.      ], FalsePositiveRate-[0.      0.      0.02 0.056
0.12 0.248 0.428 0.668 1.      ]Threshold-[2.      1.      0.85714286
0.71428571 0.57142857 0.42857143
0.28571429 0.14285714 0.      ]

```

[75]: Text(0, 0.5, 'TruePositiveRate')




```
[80]: #PrecisionRecallCurveforLogisticRegression

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs=model.predict_proba(features)
#keep probabilities for the positive outcome only
probs=probs[:,1]
# predict class values
yhat=model.predict(features)
# calculate precision-recall curve
precision,recall,thresholds=precision_recall_curve(label,probs)
#calculate F1 score
f1=f1_score(label,yhat)
# calculate precision-recall AUC
auc=auc(recall,precision)
#calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3fauc=%.3fap=%.3f'%(f1,auc,ap))#plotnoskill
plt.plot([0,1],[0.5,0.5],linestyle='--')
#plot the precision-recall curve for the model
plt.plot(recall,precision,marker='.')
```

f1=0.624auc=0.726ap=0.727

[80]: [



```
[82]: #PrecisionRecallCurveforKNNAlgorithms


from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities
probs=model2.predict_proba(features)
#keep probabilities for the positive outcome only
probs=probs[:,1]
# predict class values
yhat=model2.predict(features)
# calculate precision-recall curve
precision,recall,thresholds=precision_recall_curve(label,probs)
#calculate F1 score
f1=f1_score(label,yhat)
# calculate precision-recall AUC
auc=auc(recall,precision)
#calculate average precision score
ap = average_precision_score(label, probs)
print('f1=%.3fauc=%.3fap=%.3f'%(f1, auc, ap))
#plotnoskill
```



```
plt.plot([0,1],[0.5,0.5],linestyle='--')
#plot the precision-recall curve for the model
plt.plot(recall,precision,marker='.')
```

f1=0.658auc=0.752ap=0.709

[82]: [<matplotlib.lines.Line2D at 0x7f9c97cf4d50>]



```
[83]: #PrecisionRecallCurveforDecissionTreeClassifier

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score

# predict probabilities
probs=model3.predict_proba(features)

#keep probabilities for the positive outcome only
probs=probs[:,1]

# predict class values
#yhat = model3.predict(features)

# calculate precision-recall curve
#precision,recall,thresholds=precision_recall_curve(label, probs)
```

```

#calculateF1 score
f1=f1_score(label,yhat)
# calculate precision-recall AUC
auc=auc(recall,precision)

#calculateaverage precision score
ap = average_precision_score(label, probs)
print('f1=%.3fauc=%.3fap=%.3f'%(f1,auc,ap))


#plotnoskill
plt.plot([0,1],[0.5,0.5],linestyle='--')

#plotthe precision-recall curve for the model
plt.plot(recall,precision,marker='.')

```

f1=0.658auc=0.752ap=0.765

[83]: [<matplotlib.lines.Line2D at 0x7f9c97c3ee10>]



```

[84]: #PrecisionRecallCurveforRandomForest

from sklearn.metrics import precision_recall_curve
from sklearn.metrics import f1_score
from sklearn.metrics import auc
from sklearn.metrics import average_precision_score
# predict probabilities

```


```

probs=model4.predict_proba(features)
#keep probabilities for the positive outcome only
probs=probs[:,1]
#predict class values
yhat=model4.predict(features)
# calculate precision-recall curve
precision,recall,thresholds=precision_recall_curve(label,probs)
#calculateF1 score
f1=f1_score(label,yhat)
# calculate precision-recall AUC
auc=auc(recall,precision)
#calculateaverage precision score
ap = average_precision_score(label, probs)
print('f1=%.3fauc=%.3fap=%.3f'%(f1,auc,ap))#plotnoskill
plt.plot([0,1],[0.5,0.5],linestyle='--')
#plotthe precision-recall curve for the model
plt.plot(recall,precision,marker='.')

```

f1=0.759auc=0.802ap=0.663

[84]: [<matplotlib.lines.Line2D at 0x7f9c8fc98150>]



3.0.1 THANKYOU

[]: